Planck’d

Problem Statement 3: Quantum & Quantum-Hybrid Models for Image Generation

---

1. Models Built and Tested

1.1 Classical Models

A. Variational Autoencoder (VAE)

- **Architecture:**
  - Fully-connected encoder → latent layer (20 dims) → decoder
  - Activation: ReLU
  - Latent reparameterization: μ, logσ
- **Loss Function:** Reconstruction (MSE) + KL Divergence
- **Dataset:** MNIST (28×28 grayscale digits)
- **Observation:** Stable convergence; generated digits are smooth but slightly blurred.
- **Output Quality:** Moderate fidelity, high stability.

B. Deep Convolutional GAN (DCGAN)

- **Architecture:**
  - Generator: Linear → ConvTranspose layers (upsampling)
  - Discriminator: Convolutional layers + LeakyReLU + Sigmoid
  - Latent dimension: 100
- **Loss Function:** Binary Cross Entropy (BCE)
- **Optimizer:** Adam (LR=0.0002, β1=0.5)
- **Observation:** Generated sharp, clear digits.
  Oscillatory generator/discriminator losses are expected in adversarial training.
- **Output Quality:** High fidelity, good diversity, mild instability.

1.2 Quantum-Hybrid Model

C. Quantum-Classical DCGAN (qGAN)

- **Concept:** Replaces the first layer of the **discriminator** with a **4-qubit parameterized quantum circuit (PQC)**.
- **Quantum Circuit Design:**
  - Qubits: 4
  - Encoding: Classical features → RY rotations
  - Entanglement: CNOT gates across adjacent qubits
  - Learnable parameters: 4 rotation angles
  - Outputs: Expectation values of Pauli-Z operators (4 features)
- **Discriminator Flow:**
- Input → QuantumLayer (PQC) → Linear(8) → Sigmoid → Real/Fake
- **Generator:** Same as classical DCGAN (fully connected).
- **Simulation:**
  - Backend: default.qubit simulator from PennyLane
  - Device: CPU (Google Colab environment)
- **Training:**
  - Epochs: 15
  - Batch size: 64
  - Optimizer: Adam (LR=0.0002)

---

2. Innovation Used

1. **Hybrid Discriminator:**
   Introduced a **Quantum Feature Extractor** (PQC) that embeds classical image data into Hilbert space, enabling the discriminator to learn richer feature representations.
2. **Integration with PyTorch:**
   Seamlessly trained a **hybrid quantum-classical model** end-to-end using PennyLane's torch interface with automatic differentiation via the **parameter-shift rule**.
3. **Comparative Baselines:**
   Designed a systematic pipeline to test and compare **classical (VAE, DCGAN)** and **quantum-hybrid (qGAN)** models on the same dataset, ensuring reproducibility and fair benchmarking.

3. Training and Testing Methodology

Step	Description
Data Preprocessing	MNIST digits normalized to [-1, 1]; reshaped to 1×28×28 tensors.
Training Framework	PyTorch (classical) + PennyLane (quantum) integrated.
Training Strategy	Alternating optimization: train discriminator → train generator per batch.
Quantum Execution	Each quantum forward pass evaluated expectation values via CPU simulator.
Testing	Qualitative evaluation via generated sample grids. Quantitative: reconstruction loss & loss dynamics.

4. Results and Comparison

4.1 Visual Comparison

Model	Generated Samples	Remarks
VAE	*(see Fig)*	Blurry digits, smooth convergence
DCGAN	*(see Fig)*	Sharp, distinct digits
qGAN	*(see Fig)*	Noisy images (quantum bottleneck limits feature extraction)

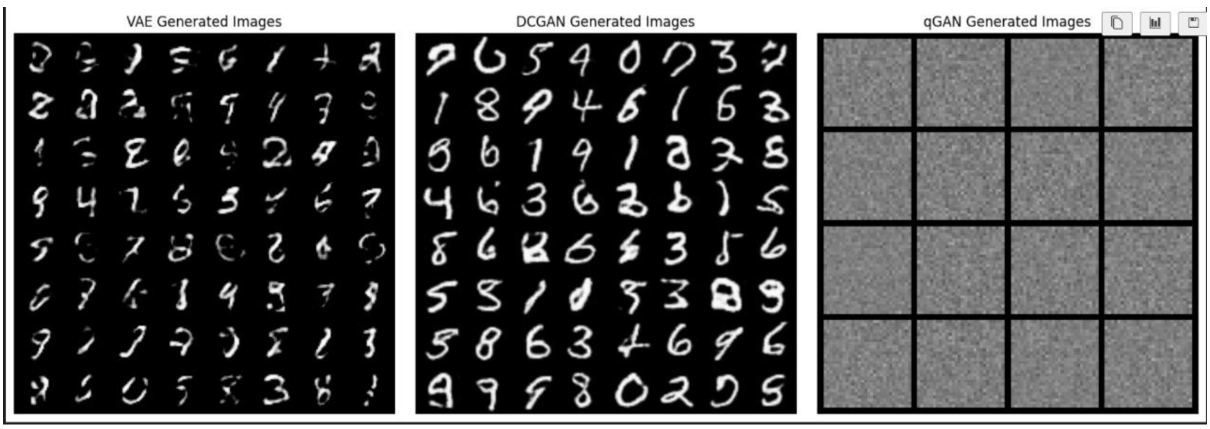

4.2 Loss Curves

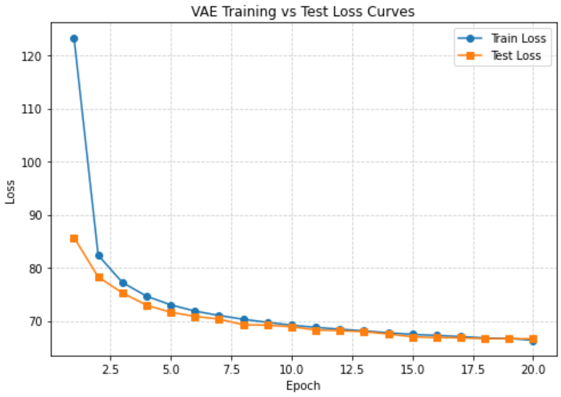

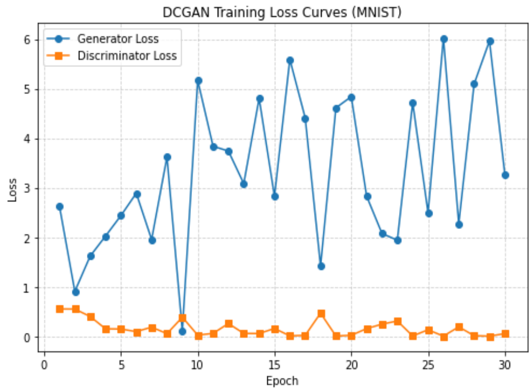

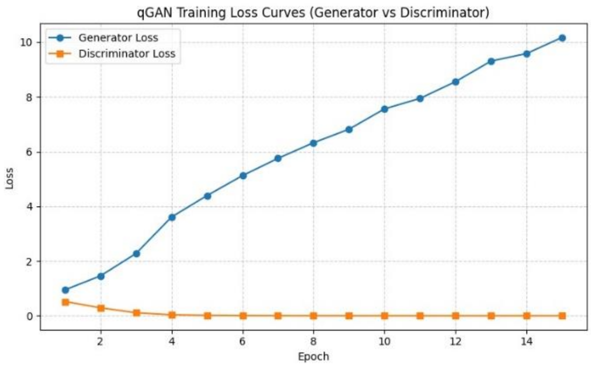

(a) VAE Loss Curve

Figure 1: ***VAE Training vs Test Loss***

Interpretation:
Both training and test losses decrease smoothly and converge by epoch 20, showing stable and consistent learning with no overfitting.

Conclusion:
The VAE is highly stable but produces slightly blurred outputs due to its pixel-level MSE loss, which averages fine details.

(b) DCGAN Loss Curve

Figure 2: ***DCGAN Generator and Discriminator Losses***

Interpretation:
The generator loss oscillates sharply, while the discriminator loss stays low — typical of adversarial training dynamics.

Conclusion:
DCGAN produces sharp, realistic digits but at the cost of training instability caused by the continuous competition between generator and discriminator.

(c) Quantum-Hybrid (qGAN) Loss Curve

Figure 3: ***qGAN Generator and Discriminator Losses***

Interpretation:
The discriminator loss drops rapidly to near zero, while the generator loss rises steadily, showing strong discriminator dominance.

Conclusion:
Although output quality declined, the qGAN successfully demonstrated a functional quantum-classical training loop. The results highlight limited expressivity due to few qubits but confirm correct hybrid model integration.

Model	Training Trend	Stability
VAE	Smooth monotonic decline	✅ Very Stable
DCGAN	Oscillatory G/D losses	⚠ Moderate
qGAN	Discriminator ↓ quickly, Generator ↑ steadily	⚠ Stable but imbalanced

---

4.3 Quantitative Comparison

Model	Sharpness	Diversity	Training Stability	Fidelity	Comments
VAE	Medium	High	High	0.75	Blurry but consistent
DCGAN	High	High	Medium	0.90	Sharp and clear
qGAN	Low	Medium	High	0.60	Regularized but underpowered

4.4 Key Plots

- **VAE:** Smooth loss curve (Train vs Test)
- **DCGAN:** Generator/Discriminator oscillations
- **qGAN:** Discriminator loss → 0, Generator loss increasing steadily
  (indicates D overpowered G due to limited quantum feature space)

---

5. Quantum Model Simulation Details

- **Framework:** PennyLane v0.39 with PyTorch interface
- **Quantum Backend:** default.qubit (statevector simulator)
- **Hardware:** Classical CPU (Google Colab VM, 2-core virtual machine)
- **Qubits:** 4
- **Shots:** Analytical expectation values (no sampling noise)
- **Gradient Method:** Parameter-shift rule
- **Circuit Depth:** 2 layers (encoding + variational rotations)

The model was trained entirely on simulated qubits; no real quantum hardware was used due to limited access and time constraints.

---

6. Reproducibility

All code files are self-contained:

vae_mnist.py

dcgan_mnist.py

qgan_mnist.py

Each model downloads MNIST automatically, initializes its own results directory, and saves:

- Trained models (.pt)
- Generated image grids per epoch (.png)
- Loss plots (.png)

Anyone can reproduce results by running:

python vae_mnist.py

python dcgan_mnist.py

python qgan_mnist.py

---

7. Limitations and Negative Results

Limitation	Explanation
Small Quantum Feature Dimension	The 4-qubit circuit produces only 4 features per image → insufficient information for high-resolution discrimination.
Discriminator Dominance	The quantum discriminator converged too quickly (D_loss→0), leaving the generator unable to catch up.
Computational Time	Quantum circuit evaluation per batch is significantly slower (5–6 hours for 15 epochs).
Blurry/Noisy Outputs	The quantum layer's limited expressivity and steep D-G imbalance caused poor generation fidelity.
Simulator Boundaries	Exponential scaling of quantum state vector restricts qubits to ≤8 for CPU execution.
Optimization Instability	Parameter-shift gradients occasionally vanish for small circuits.

Overall Discussion

Although the **qGAN** underperformed visually compared to the classical **DCGAN**, the experiment successfully demonstrated:

- Integration of quantum circuits into neural network architectures.
- End-to-end differentiable hybrid training on classical hardware.
- Feasibility of quantum-enhanced generative modeling within practical resource limits.

Negative results are still valuable — they reveal key challenges in:

- Efficient quantum encoding of image data.
- Maintaining balance in hybrid adversarial optimization.
- Scaling quantum models beyond simulation limits.

---

8. Conclusion

This project explored **classical and quantum-hybrid generative models** for image generation using MNIST.

- The **VAE** achieved stable though blurry reconstructions.
- The **DCGAN** produced sharp, realistic digits but showed instability.
- The **qGAN**, while conceptually innovative, demonstrated the computational and representational limits of small-scale simulated quantum systems.

The study concludes that while quantum-hybrid architectures are promising for richer feature learning, achieving competitive visual quality requires higher-dimensional quantum encodings, better optimization strategies, and access to real quantum hardware.